



Learn the AI/ML Project on:
Iris-Species Classification



Find all projects at @ github.com/Masterx-AI

★ AI / ML Project - Iris Species Classification ★



Description:

The Iris dataset was used in R.A. Fisher's classic 1936 paper, The Use of Multiple Measurements in Taxonomic Problems, and can also be found on the UCI Machine Learning Repository.

It includes three iris species with 50 samples each as well as some properties about each flower. One flower species is linearly separable from the other two, but the other two are not linearly separable from each other.

The columns in this dataset are:

- Id
- SepalLengthCm
- SepalWidthCm
- PetalLengthCm
- PetalWidthCm
- Species

Objective:

- Import the Iris Dataset from sklearn library.
- Build classification models to predict the species.
- Compare the evaluation metrics of various classification algorithms.

1. Data Exploration

In [1]:

```
#Importing the basic librarieres

import numpy as np
import pandas as pd
import seaborn as sns
from IPython.display import display

import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = [10,6]

import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
#Importing the dataset

from sklearn.datasets import load_iris

Iris_Dataset = load_iris()
labels = load_iris().target_names
data = pd.DataFrame(load_iris().data, columns=['Sepel_Length', 'Sepel_Width', 'Petal_Length',
target = pd.DataFrame(load_iris().target, columns=['Species']))
df = pd.concat([data,target], axis=1)
display(df.head(5))
print('\n\033[1mInference:\033[0m The Datset consists of {} features & {} samples.'.format(
```

	Sepel_Length	Sepel_Width	Petal_Length	Petal_Width	Species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

Inference: The Datset consists of 5 features & 150 samples.

In [3]:

```
#Checking the dtypes of all the columns

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sepal_Length    150 non-null    float64
1   Sepal_Width     150 non-null    float64
2   Petal_Length    150 non-null    float64
3   Petal_Width     150 non-null    float64
4   Species         150 non-null    int32
dtypes: float64(4), int32(1)
memory usage: 5.4 KB
```

In [4]:

```
#Checking the stats of all the columns

display(df.describe())
print('\n \033[1mInference:\033[0m The stats seem to be fine, let us do further analysis on
```

	Sepal_Length	Sepal_Width	Petal_Length	Petal_Width	Species
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333	1.000000
std	0.828066	0.435866	1.765298	0.762238	0.819232
min	4.300000	2.000000	1.000000	0.100000	0.000000
25%	5.100000	2.800000	1.600000	0.300000	0.000000
50%	5.800000	3.000000	4.350000	1.300000	1.000000
75%	6.400000	3.300000	5.100000	1.800000	2.000000
max	7.900000	4.400000	6.900000	2.500000	2.000000

Inference: The stats seem to be fine, let us do further analysis on the Dataset

2. Exploratory Data Analysis (EDA)


In [5]:

```
#Let us first analyze the distribution of the target variable
```

```
print('\033[1mTarget Variable Distribution'.center(55))
plt.pie(df.Species.value_counts(), labels=labels, counterclock=False, shadow=True, explode=
plt.show()
```

```
print('\n\033[1mInference:\033[0m The Target Variable seems to be perfectly balanced!')
```

Target Variable Distribution



Inference: The Target Variable seems to be perfectly balanced!

In [6]:

```
#Understanding the feature set
```

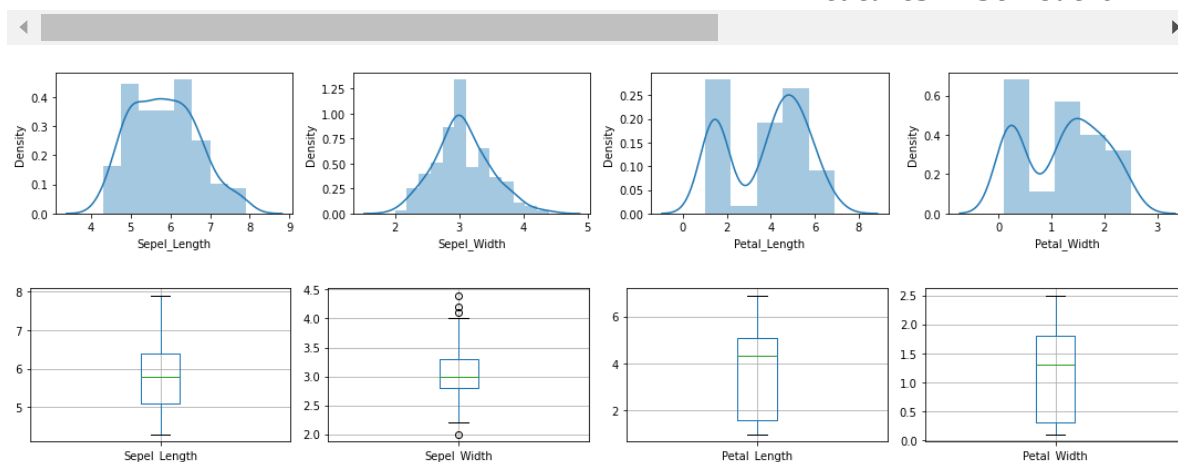
```
print('\033[1mFeatures Distribution'.center(130))
```

```
plt.figure(figsize=[15,2.5])
for c in range(4):
    plt.subplot(1,4,c+1)
    sns.distplot(df[df.columns[c]])
plt.tight_layout()
plt.show()
```

```
plt.figure(figsize=[15,2.5])
for c in range(4):
    plt.subplot(1,4,c+1)
    df.boxplot(df.columns[c])
plt.tight_layout()
plt.show()
```

```
print('\n\033[1mInference:\033[0m The distribution of Petal Length & Petal Width is not normal. Also there are some outliers in the Sepal Width. Let us fix them in the upcoming section')
```

Features Distribution



Inference: The distribution of Petal Length & Petal Width is not normal. Also there are some outliers in the Sepal Width. Let us fix them in the upcoming section

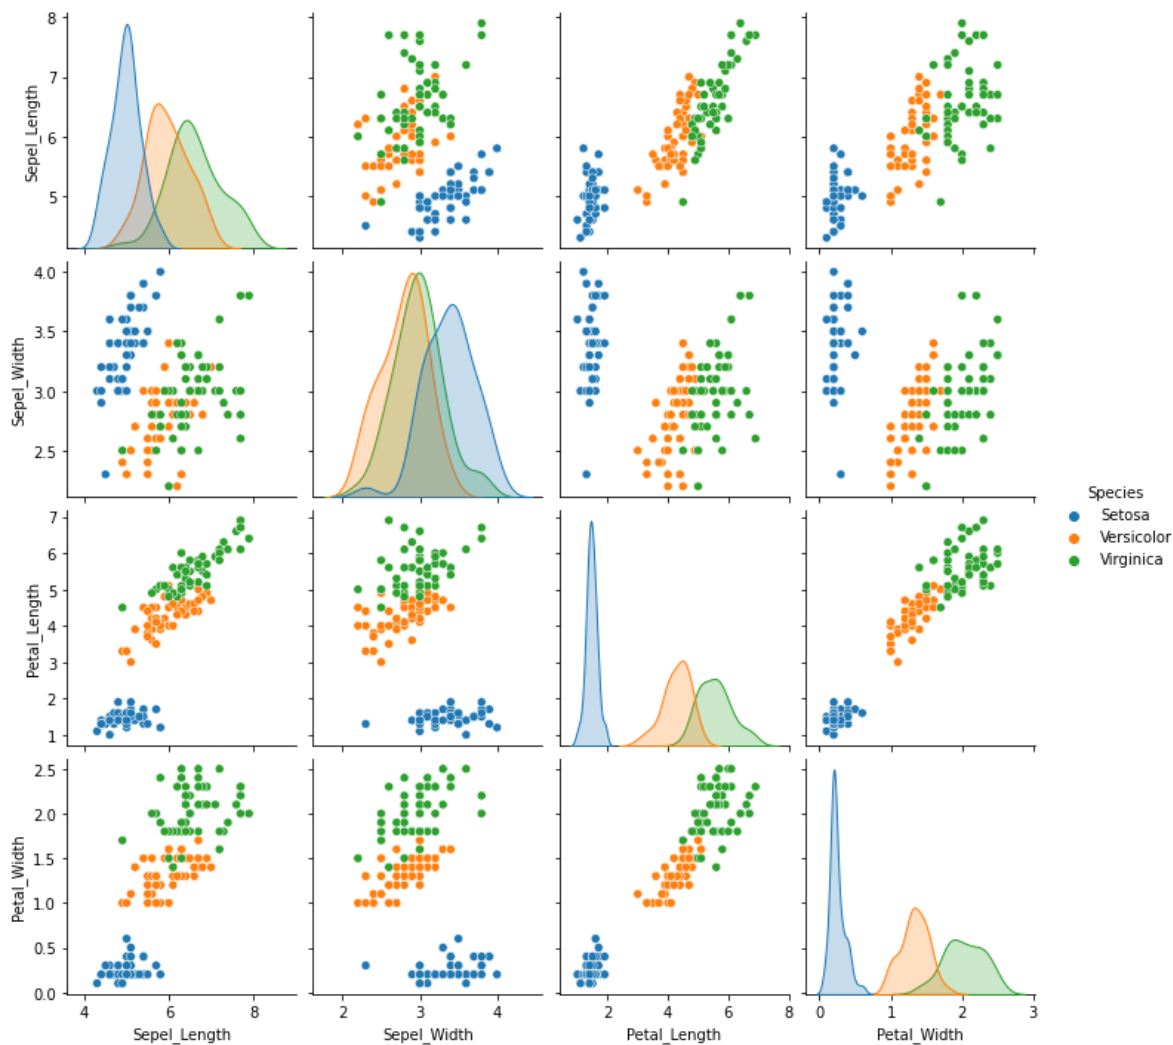
In [214]:

```
#Understanding the relationship between all the features
```

```
df1 = df.copy()
df1.Species = df.Species.map({0:'Setosa',1:'Versicolor',2:'Virginica'})
```

```
sns.pairplot(df1, hue="Species")
plt.show()
```

```
print('\n\033[1mInference:\033[0m We Observe that the Setosa can be clearly distinguished,
some overlap between Versicolor & Virginica')
```



Inference: We Observe that the Setosa can be clearly distinguished, while there is some overlap between Versicolor & Virginica

3. Data Preprocessing

In [8]:

```
#Check for empty elements
```

```
print(df.isnull().sum())  
print('\n\033[1mInference:\033[0m The dataset doesn\'t have any null elements')
```

```
Sepal_Length    0  
Sepal_Width     0  
Petal_Length    0  
Petal_Width     0  
Species         0  
dtype: int64
```

Inference: The dataset doesn't have any null elements

In [9]:

```
#Removal of any Duplicate rows (if any)
```

```
counter = 0  
r,c = df1.shape  
  
df.drop_duplicates(inplace=True)  
  
if df.shape==(r,c):  
    print('\n\033[1mInference:\033[0m The dataset doesn\'t have any duplicates')  
else:  
    print(f'\n\033[1mInference:\033[0m Number of duplicates dropped/fixed ---> {r-df.shape[0]}')
```

Inference: Number of duplicates dropped/fixed ---> 1

In [10]:

```
#Removal of outlier:

for i in df.columns:
    Q1 = df[i].quantile(0.25)
    Q3 = df[i].quantile(0.75)
    IQR = Q3 - Q1
    df = df[df[i] <= (Q3+(1.5*IQR))]
    df = df[df[i] >= (Q1-(1.5*IQR))]
    df = df.reset_index(drop=True)
display(df)
print('\n\033[1mInference:\033[0m After removal of outliers, The dataset now has {} feature
```

	Sepal_Length	Sepal_Width	Petal_Length	Petal_Width	Species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0
...
140	6.7	3.0	5.2	2.3	2
141	6.3	2.5	5.0	1.9	2
142	6.5	3.0	5.2	2.0	2
143	6.2	3.4	5.4	2.3	2
144	5.9	3.0	5.1	1.8	2


145 rows × 5 columns

Inference: After removal of outliers, The dataset now has 5 features & 145 samples.

In [11]:

```
#Final Dataset size after performing Preprocessing
```

```
plt.title('Final Dataset Samples')
plt.pie([df.shape[0], df1.shape[0]-df.shape[0]], radius = 1, labels=['Retained', 'Dropped'],
        autopct='%1.1f%%', pctdistance=0.9, explode=[0,0], shadow=True)
plt.pie([df.shape[0]], labels=['100%'], labeldistance=-0, radius=0.78)
plt.show()
```



4. Feature Scaling

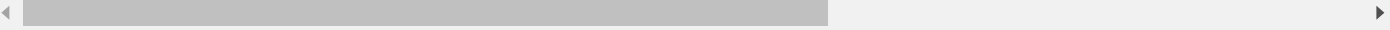
In [12]:

```
#Splitting the data into training & testing sets

from sklearn.model_selection import train_test_split

X = df.drop(['Species'],axis=1)
Y = df.Species
Train_X, Test_X, Train_Y, Test_Y = train_test_split(X, Y, train_size=0.8, test_size=0.2, ra

print('Original set ---> ',X.shape,Y.shape,'\nTraining set ---> ',Train_X.shape,Train_Y.s
```



```
Original set ---> (145, 4) (145,)
Training set ---> (116, 4) (116,)
Testing set ---> (29, 4) (29,)
```

In [13]:

```
#Feature Scaling (Standardization)

from sklearn.preprocessing import StandardScaler

std = StandardScaler()

Train_X_std = std.fit_transform(Train_X)
Test_X_std = std.transform(Test_X)
```

5. Predictive Modeling

In [145]:

```
#Let us create first create a table to store the results of various models

Evaluation_Results = pd.DataFrame(np.zeros((8,4)), columns=['Accuracy', 'Precision', 'Recall', 'F1-score'])
Evaluation_Results.index=['Logistic Regression (LR)', 'Decision Tree Classifier (DT)', 'Random Forest Classifier (RF)', 'Support Vector Machine (SV)', 'K Nearest Neighbours (KNN)', 'Gradient Boosting (GB)', 'Extreme Gradient Boosting (XGB)']

Evaluation_Results
```

Out[145]:

	Accuracy	Precision	Recall	F1-score
Logistic Regression (LR)	0.0	0.0	0.0	0.0
Decision Tree Classifier (DT)	0.0	0.0	0.0	0.0
Random Forest Classifier (RF)	0.0	0.0	0.0	0.0
Naïve Bayes Classifier (NB)	0.0	0.0	0.0	0.0
Support Vector Machine (SV)	0.0	0.0	0.0	0.0
K Nearest Neighbours (KNN)	0.0	0.0	0.0	0.0
Gradient Boosting (GB)	0.0	0.0	0.0	0.0
Extreme Gradient Boosting (XGB)	0.0	0.0	0.0	0.0

In [219]:

```
#Let us define functions to summarise the Prediction's scores .
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, f1_score
```

```
#Classification Summary Function
```

```
def Classification_Summary(pred,i):
    Evaluation_Results.iloc[i]['Accuracy']=round(accuracy_score(Test_Y, pred),3)*100
    Evaluation_Results.iloc[i]['Precision']=round(precision_score(Test_Y, pred, average='weighted'),3)*100
    Evaluation_Results.iloc[i]['Recall']=round(recall_score(Test_Y, pred, average='weighted'),3)*100
    Evaluation_Results.iloc[i]['F1-score']=round(f1_score(Test_Y, pred, average='weighted'),3)*100
    print('{}{}\033[1m{}\033[0m{}\n'.format('< '*3,'- '*35,Evaluation_Results.index[i], '-'
    print('Accuracy = {}'.format(round(accuracy_score(Test_Y, pred),3)*100))
    print('F1 Score = {}'.format(round(f1_score(Test_Y, pred, average='weighted'),3)*100))
    print('\n \033[1mConfusion Matrix:\033[0m\n',confusion_matrix(Test_Y, pred))
    print('\n \033[1mClassification Report:\033[0m\n',classification_report(Test_Y, pred))
```

```
#Visualising Function
```

```
def ROC_plot(m):
    pred = m.predict(Test_X)
    ref = [0 for _ in range(len(Test_Y))]
    ref_auc = roc_auc_score(Test_Y, ref)
    lr_auc = roc_auc_score(Test_Y, pred)

    ns_fpr, ns_tpr, _ = roc_curve(Test_Y, ref)
    lr_fpr, lr_tpr, _ = roc_curve(Test_Y, pred)

    plt.plot(ns_fpr, ns_tpr, linestyle='--')
    plt.plot(lr_fpr, lr_tpr, marker='.', label='Area = {}'.format(round(roc_auc_score(Test_Y, pred),3)*100))
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.legend()
    plt.show()
```

In [176]:

```
#Logistic Regression
```

```
from sklearn.linear_model import LogisticRegression
```

```
LR = LogisticRegression().fit(Train_X_std, Train_Y)
```

```
pred = LR.predict(Test_X_std)
```

```
Classification_Summary(pred,0)
```

```
<<<-----Logistic Regression (LR)-----  
----->>>
```

Accuracy = 89.7%

F1 Score = 89.9%

Confusiton Matrix:

```
[[ 8  0  0]
```

```
 [ 0  5  1]
```

```
 [ 0  2 13]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8
1	0.71	0.83	0.77	6
2	0.93	0.87	0.90	15
accuracy			0.90	29
macro avg	0.88	0.90	0.89	29
weighted avg	0.90	0.90	0.90	29

In [177]:

```
from sklearn.tree import DecisionTreeClassifier
DT = DecisionTreeClassifier().fit(Train_X_std, Train_Y)
pred = DT.predict(Test_X_std)
Classification_Summary(pred,1)
```

<<<-----Decision Tree Classifier (DT)-----
----->>>

Accuracy = 96.6%
F1 Score = 96.5%

Confusiton Matrix:

```
[[ 8  0  0]
 [ 0  5  1]
 [ 0  0 15]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8
1	1.00	0.83	0.91	6
2	0.94	1.00	0.97	15
accuracy			0.97	29
macro avg	0.98	0.94	0.96	29
weighted avg	0.97	0.97	0.96	29

In [178]:

```
from sklearn.ensemble import RandomForestClassifier
RF = RandomForestClassifier().fit(Train_X_std, Train_Y)
pred = RF.predict(Test_X_std)
Classification_Summary(pred,2)
```

<<<-----Random Forest Classifier (RF)-----
----->>>

Accuracy = 89.7%
F1 Score = 89.9%

Confusiton Matrix:

```
[[ 8  0  0]
 [ 0  5  1]
 [ 0  2 13]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8
1	0.71	0.83	0.77	6
2	0.93	0.87	0.90	15
accuracy			0.90	29
macro avg	0.88	0.90	0.89	29
weighted avg	0.90	0.90	0.90	29

In [179]:

```
from sklearn.naive_bayes import GaussianNB
NB = GaussianNB().fit(Train_X_std, Train_Y)
pred = NB.predict(Test_X_std)
Classification_Summary(pred,3)
```

<<<-----Naïve Bayes Classifier (NB)-----
----->>>

Accuracy = 93.10000000000001%
F1 Score = 93.10000000000001%

Confusiton Matrix:

```
[[ 8  0  0]
 [ 0  5  1]
 [ 0  1 14]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8
1	0.83	0.83	0.83	6
2	0.93	0.93	0.93	15
accuracy			0.93	29
macro avg	0.92	0.92	0.92	29
weighted avg	0.93	0.93	0.93	29

In [180]:

```
from sklearn.svm import LinearSVC
SV = LinearSVC().fit(Train_X_std, Train_Y)
pred = SV.predict(Test_X_std)
Classification_Summary(pred,4)
```

<<<-----Support Vector Machine (SV)-----
----->>>

Accuracy = 93.10000000000001%
F1 Score = 93.10000000000001%

Confusiton Matrix:

```
[[ 8  0  0]
 [ 0  5  1]
 [ 0  1 14]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8
1	0.83	0.83	0.83	6
2	0.93	0.93	0.93	15
accuracy			0.93	29
macro avg	0.92	0.92	0.92	29
weighted avg	0.93	0.93	0.93	29

In [181]:

```
from sklearn.neighbors import KNeighborsClassifier
KNN = KNeighborsClassifier().fit(Train_X_std, Train_Y)
pred = KNN.predict(Test_X_std)
Classification_Summary(pred,5)
```

<<<-----K Nearest Neighbours (KNN)-----
----->>>

Accuracy = 93.10000000000001%
F1 Score = 93.30000000000001%

Confusiton Matrix:

```
[[ 8  0  0]
 [ 0  6  0]
 [ 0  2 13]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8
1	0.75	1.00	0.86	6
2	1.00	0.87	0.93	15
accuracy			0.93	29
macro avg	0.92	0.96	0.93	29
weighted avg	0.95	0.93	0.93	29

In [182]:

```
from sklearn.ensemble import GradientBoostingClassifier
GB = GradientBoostingClassifier().fit(Train_X_std, Train_Y)
pred = GB.predict(Test_X_std)
Classification_Summary(pred,6)
```

<<<-----Gradient Boosting (GB)-----
----->>>

Accuracy = 93.10000000000001%
F1 Score = 93.10000000000001%

Confusiton Matrix:

```
[[ 8  0  0]
 [ 0  5  1]
 [ 0  1 14]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8
1	0.83	0.83	0.83	6
2	0.93	0.93	0.93	15
accuracy			0.93	29
macro avg	0.92	0.92	0.92	29
weighted avg	0.93	0.93	0.93	29

In [183]:

```
from xgboost import XGBClassifier
XGB = XGBClassifier().fit(Train_X_std, Train_Y)
pred = XGB.predict(Test_X_std)
Classification_Summary(pred,7)
```

```
[19:56:35] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_
1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation
metric used with the objective 'multi:softprob' was changed from 'merror' to
'mlogloss'. Explicitly set eval_metric if you'd like to restore the old beha
vior.
```

```
<<<-----Extreme Gradient Boosting (XGB)-----
----->>>
```

Accuracy = 89.7%

F1 Score = 89.9%

Confusiton Matrix:

```
[[ 8  0  0]
 [ 0  5  1]
 [ 0  2 13]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8
1	0.71	0.83	0.77	6
2	0.93	0.87	0.90	15
accuracy			0.90	29
macro avg	0.88	0.90	0.89	29
weighted avg	0.90	0.90	0.90	29

In [212]:


#Plotting Confusion-Matrix of all the predictive Models

```
def plot_cm(y_true, y_pred):
    cm = confusion_matrix(y_true, y_pred, labels=np.unique(y_true))
    cm_sum = np.sum(cm, axis=1, keepdims=True)
    cm_perc = cm / cm_sum.astype(float) * 100
    annot = np.empty_like(cm).astype(str)
    nrows, ncols = cm.shape
    for i in range(nrows):
        for j in range(ncols):
            c = cm[i, j]
            p = cm_perc[i, j]
            if i == j:
                s = cm_sum[i]
                annot[i, j] = '%.1f%%\n%d/%d' % (p, c, s)
            elif c == 0:
                annot[i, j] = ''
            else:
                annot[i, j] = '%.1f%%\n%d' % (p, c)
    cm = pd.DataFrame(cm, index=np.unique(y_true), columns=np.unique(y_true))
    cm.columns=labels
    cm.index=labels
    cm.index.name = 'Actual'
    cm.columns.name = 'Predicted'
    #fig, ax = plt.subplots()
    sns.heatmap(cm, annot=annot, fmt='')# cmap= "GnBu"

def conf_mat_plot(all_models):
    plt.figure(figsize=[20,7])

    for i in range(len(all_models)):
        plt.subplot(2,4,i+1)
        pred = all_models[i].predict(Test_X_std)
        plot_cm(Test_Y, pred)
        plt.title(Evaluation_Results.index[i])
    plt.tight_layout()
    plt.show()

conf_mat_plot([LR,DT,RF,NB,SV,KNN,GB,XGB])
```




In [211]:

```
#Comparing all the models Scores
```

```
#plt.figure(figsize=[12,5])
```

```
sns.heatmap(Evaluation_Results, annot=True, vmin=80.0, vmax=100.0, cmap='YlGnBu', fmt='.1f')  
plt.show()
```



6. Project Outcomes & Conclusions

Here are some of the key outcomes of the project:

- The Dataset was quite small totally just 150 samples & after preprocessing 3.3% of the data samples were dropped. It was also balanced & didn't require any artificial techniques to balance it.
- Visualising the distribution of data & their relationships, helped us to get some insights on the class separability.
- Feature selection or feature extracting as there were only 4 features, which all contributed towards the right prediction.
- Testing multiple algorithms with default hyperparameters gave us some understanding for various models performance on this specific dataset.
- While, Decision Tree Classifier algorithm gave the best overall scores for the current dataset, yet it is wise to also consider simpler models as they are more generalisable.

In []:

<<<-----THE END-----